

Content, Citations & Indexing.

The research behind three big upgrades — bulk one-click pages, real citation robots, and search-engine indexing — and exactly what we build next, in plain language.

Research findings

What's already built

The build plan

Claude Code → Dashboard

WHAT YOU ASKED FOR

Four upgrades, one honest plan.

You gave us three big goals. The good news from the research: **a lot of the hard machinery already exists** in AIOS. This document explains each goal in simple terms, shows what is already built, and lays out what we build next — everything **drivable from Claude Code first, then the dashboard**.

1

Content — research-first bulk pages, published & editable

2

Citations — real robots across 40–50+ directories

3

Citation audit — right directories for the client

4

Indexing — brand-new module to get pages found

THE HEADLINE FINDING

✓ Already built

The content pipeline, the WordPress plugin, the **Playwright citation robot**, a **155-directory catalog**, the CAPTCHA solver, and Web 2.0 publishing to 17 platforms.

🟡 Partly there

Directory **prioritisation logic** exists; the citation robot covers **12 of 155** directories; pages publish but as **plain HTML**, not Elementor-editable.

+ To build

The **Research module**, **bulk one-click** generation, **account creation**, extra business fields, the **audit view**, and the whole **Indexing module**.

HOW TO READ THIS


Green means it works today. Amber means the foundation is laid but needs finishing. Maroon means it is genuinely new work. We were careful not to call anything "done" that isn't — the last two pages give you the full scorecard and the phased timeline.

THE VISION — ONE CLICK, MANY PAGES

From "write me a page" to "build my whole site."

A business may need 10, 15, even 20 service pages. Today we make them one at a time. The upgrade: you pick what you want, the system **researches, plans, writes, publishes and hands back the live URLs** — for one page or all of them at once.

THE AUTOMATIC PIPELINE (RUNS END-TO-END ONCE YOU HIT GO)



THE RESEARCH MODULE — DECIDE BEFORE YOU WRITE

Tell it the type. It finds the pages.

Before any writing happens, a dedicated **Research module** asks what kind of pages you want, studies the website, and comes back with a recommended list you simply tick.

PICK A CONTENT TYPE

Service pages

Location pages

Service + Location

Service-area pages

GSC / opportunity pages

Blogs

FAQs

...more

WHAT IT RECOMMENDS — YOU TICK, IT BUILDS

DBC.COM · SERVICE PAGES · 15 RECOMMENDED

☒ Emergency Plumbing 1,900 vol · easy

☒ Drain Cleaning 1,300 vol · easy

☒ Water Heater Repair 880 vol · med

☐ Leak Detection 720 vol · easy

☒ Bathroom Remodeling 2,400 vol · hard

Generate selected (4)

Generate all 15

SMART FOR EVERY TYPE

Service + Location

Researches every sensible **city × service** combination — "Emergency Plumbing in Parramatta", "Drain Cleaning in Bondi" — and recommends the winners.

Location / Service-area

Maps the towns and suburbs the business actually serves and proposes a page for each.

Blogs & FAQs

Mines questions people ask and gaps competitors miss, then proposes topics.

REAL PAGES — NOT JUST TEXT

Editable pages that match the brand.

This is the most important content requirement — and today's biggest gap. Right now our plugin publishes a page as **fixed HTML**. It needs to publish pages your team can **drag, edit and restyle inside WordPress / Elementor**.

Think of it like this

Fixed HTML is a **printed poster** — it looks fine but you can't rearrange it. An Elementor page is **LEGO** — every heading, image and button is a block you can move and edit. To hand over LEGO, our plugin must give WordPress the page as **building blocks**, not a flat poster.

Design & brand consistency

If the site already has a design: new pages copy its existing Elementor style kit — same fonts, colours, buttons. **If it doesn't:** we generate a clean layout using best practice for that niche.


BEFORE → AFTER

TODAY



Plain HTML. Looks OK live, but the client can't easily edit sections in the builder.

AFTER



Editable blocks. Every section is a real Elementor widget the client can move & restyle.

THE BUILD, IN ONE LINE

Teach the plugin to write Elementor's own page format (its block data) instead of flat HTML — plus a step that reads and reuses the site's existing style. Whether we extend today's plugin or ship a new one doesn't matter; the result must be **fully editable pages**.

CONTENT — WHAT'S BUILT VS WHAT'S NEW

Most of the engine is already running.

6

core pieces already built & tested

4

new pieces to add on top

17

WordPress page-type frameworks ready

✓

publishes live as a draft (proven on spotino.org)

CAPABILITY	STATUS	NOTES
Content pipeline & workflow	BUILT	Queued → draft → review → publish, with a human approval gate.
AI writer + SEO doctrine	BUILT	Ranking-grade drafts with SERP teardown & a quality scorecard (now advisory, not a blocker).
WordPress publishing	BUILT	Three routes — REST, XML-RPC (login password), and our plugin. Lands as a draft.
Page-type frameworks	BUILT	Service, local, blog, FAQ structures auto-selected per page.
Research module (competitor + recommend)	NEW	The "study the site, propose the page set" front end.
Bulk / one-click fan-out	NEW	Generate many selected pages in one go, collect all URLs.
Elementor-editable output	NEW	Plugin writes builder blocks, not flat HTML.
Design / brand matching	NEW	Reuse the site's existing style, or niche best-practice.

WHY THE ROBOTS, NOT THE APIS

APIs were the lesson. Playwright is the goal.

The citation APIs we requested only cover a handful of platforms — they were always meant to **teach us the workflow**, not be the final answer. The real system: pick "citations" from a dropdown, click once, and a **browser robot submits across 40–100 directories** and hands back the URLs.

What APIs gave us

- Understanding of how a submission works
- The exact fields directories ask for
- Where things fail (accounts, CAPTCHAs)
- Coverage of only **a few** platforms

What Playwright gives us

- One click → **dozens** of directories
- Works on sites that have **no API at all**
- A screenshot proof for every submission
- ~2.5× cheaper than off-the-shelf actors

GOOD NEWS FROM THE CODE

The Playwright robot **already exists** in AIOS — it launches a real browser, fills a directory's form from a simple data recipe, solves CAPTCHAs, and screenshots the result as proof. A new directory is added as **data, not new code**. Today it ships with **12 directory recipes** against a catalog of **155**. The job now is coverage, accounts, and turning it on.

HOW THE CITATION ROBOT IS PUT TOGETHER

One engine, 155 recipes.

The engine never changes. Each directory is just a small **recipe** — its web address, which box takes the name, the address, the phone — plus how to tell it worked. Add a directory by adding a recipe.



THE 155-DIRECTORY CATALOG — BY HOW WE SUBMIT



- Bot-fillable — the robot can just fill & submit
- CAPTCHA-assisted — robot + a solver
- Manual-only — needs a human
- Aggregator / API

...AND BY COUNTRY




What this means: ~128 of the 155 (bot-fillable + CAPTCHA) are automatable today by writing their recipes. Reaching the "40–50 live" target is well within the first push.

THE REAL GAP — CREATING ACCOUNTS

Every directory wants an account first.

Today the robot fills "add your business" forms — but most directories make you **register and confirm an email** before that. There are no agency accounts for these sites (if there were, an API would already solve it). So the robot must **create an account per client**, and that flow is the main new build.

THE ACCOUNT-CREATION FLOW TO BUILD



Per-client accounts

Each client gets their own logins, stored safely in the vault. Cleaner, and it's how these directories expect it.

The email piece

We add an inbox the robot can read, so confirmation links get clicked automatically — the part that's missing today.

Honest scope

This is the hard, new subsystem. We build it once, prove it on a handful of sites, then scale directory by directory.

PROOF TARGET

First milestone: real citations submitted on a batch of directories using a sample business — each returning a **live submission URL and a screenshot**. It doesn't matter whose NAP we use to prove it; it matters that the listing goes up and we can show where.

SHOW THE CLIENT ONLY WHAT HELPS THEM

The right directories, in the right order.

With 155+ directories, dumping them all on a client is noise. An **Australian plumber** gains nothing from a US-only site. The audit collects the business info, works out the **country and niche**, and recommends in a clear priority order.

1 • Generic

Yelp-type — every business needs these

2 • Country-specific

Directories for the client's country

3 • Niche-specific

Directories for the client's industry

Worked example

Australian plumbing business →

- **Generic:** Yelp, Google, Bing
- **Country:** True Local, Yellow Pages AU
- **Niche:** plumbing & trade directories

US & UK-only sites are hidden.

ALREADY PARTLY BUILT

The selection engine already ranks directories by **authority tier**, filters by **country** and by **niche**, and skips low-value ones. The new work is mostly **surfacing it as an audit** — "here's built vs missing, prioritised for you" — and letting the client pick what to build from that list.

MORE THAN A NAME AND NUMBER

Directories ask for 10–15 details.

Today we mostly collect **NAP** — name, address, phone. But directories ask for much more, and a fuller listing ranks better. We add input fields for everything a directory might want.

COLLECTED TODAY

Business name	Address	City / region
Postcode	Phone	Website
Hours*	Categories	

*Hours are stored but not yet sent to the robot.

TO ADD

Business description	Logo / photos	Email
Facebook	Instagram	LinkedIn
Year founded	Payment types	
Service area	Tagline	

A SINGLE BUSINESS PROFILE POWERS EVERY SUBMISSION

BUSINESS NAME	PHONE	EMAIL NEW
DESCRIPTION NEW	LOGO NEW	
	upload	

Fill it once per client → the robot reuses it across all 40–50+ directories.

THE LAST TWO PIECES

Web 2.0 backlinks & getting found.

Web 2.0 — mostly built

Publishing to **17 platforms** (Telegra.ph, and more) already works over their APIs, returning a live URL each time.

Telegra.ph

Blogger

Medium

Tumblr


+13 more

Same as citations: where a platform has **no API**, the Playwright robot extends coverage using the same recipe approach.

Indexing — brand new

Publishing a page isn't enough — Google has to **find** it. There is **no indexing module today**. We build one that pings the search engines the moment a page goes live.

THE INDEXING FLOW



17

Web 2.0 platforms live via API today

0

indexing capability today — a clean-slate build

1-click

auto-index every page we publish

EVERYTHING, IN ONE TABLE

Built • partial • new.

No spin — here is the true state of every capability across the four workstreams.

CAPABILITY	STATE	WHAT REMAINS
● Content pipeline + AI writer + WP publish	BUILT	Nothing — proven live.
● Research module + bulk pages	NEW	Study site/competitors, recommend, fan-out.
● Elementor-editable + design match	NEW	Plugin writes builder blocks; reuse site style.
● Playwright citation robot (engine)	BUILT	Engine, CAPTCHA solver, proof screenshots done.
● Directory coverage (recipes)	12 / 155	Write recipes to reach 40–50+ live.
● Per-client account creation + email	NEW	Signup + inbox + confirm-link flow.
● Citation audit (geo / niche / generic)	PARTIAL	Logic exists; surface it as an audit view.
● Business fields beyond NAP	PARTIAL	Add ~8 fields (description, logo, socials...).
● Web 2.0 publishing (17 platforms)	BUILT	Extend to browser automation where no API.
● Indexing module	NEW	IndexNow + Google Indexing + sitemap ping.

4

Built

3

Partial

4

New

1

brand-new module
(indexing)

HOW WE BUILD IT — AND IN WHAT ORDER

Four phases. Claude Code first.

Each phase ships end-to-end and is **drivable from Claude Code first** (so we can test it hands-on), **then wired into the dashboard**. Order runs easiest-and-highest-value first, hardest-but-huge in the middle.

P1

Content

Research module + bulk pages + Elementor

Study site & competitors → recommend page set → tick → auto-write & publish as editable Elementor pages that match the brand. Builds straight on the working pipeline. **Fastest visible win.**

P2

Citations

Citation audit + expanded business fields

Surface the geo/niche/generic audit and add the extra input fields. Mostly wiring what already exists — quick, and it feeds the automation. **Low risk.**

P3

Citations

Account creation + scale to 40–50 directories

The big one — build signup + email-confirm, then write recipes across the catalog. Prove with live URLs on a first batch, then grow. **Highest effort, biggest payoff.**

P4

Reach

Web 2.0 browser automation + Indexing module

Extend the robot to no-API Web 2.0 platforms, and ship the brand-new indexing module so every published page gets found fast.

ONE INFRASTRUCTURE NOTE

The citation/Web 2.0 robots need a **browser worker** (a machine that runs Chromium) and a **client inbox** for email confirmations. These are one-time setup items we stand up in Phase 3 — flagged here so there are no surprises.